

# Trustworthy Diffusion Models & Diffusion Models for Trustworthy ML

Tianyu Pang

Sea AI Lab





# **Big Bang of Diffusion Models!**

# Diffusion Models in 2020 (Nonequilibrium Thermodynamics)

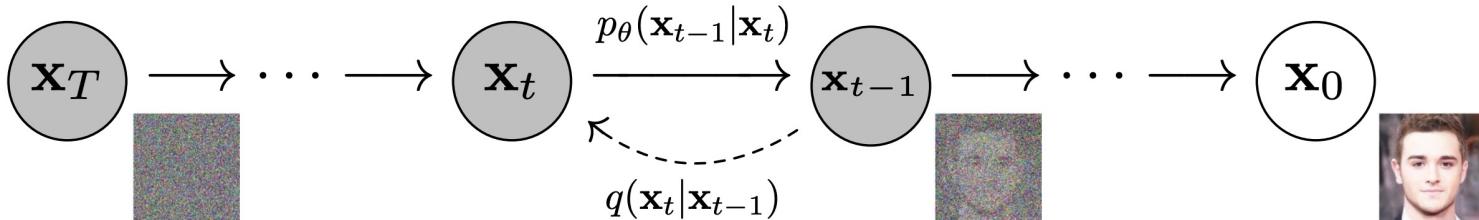


Figure 2: The directed graphical model considered in this work.

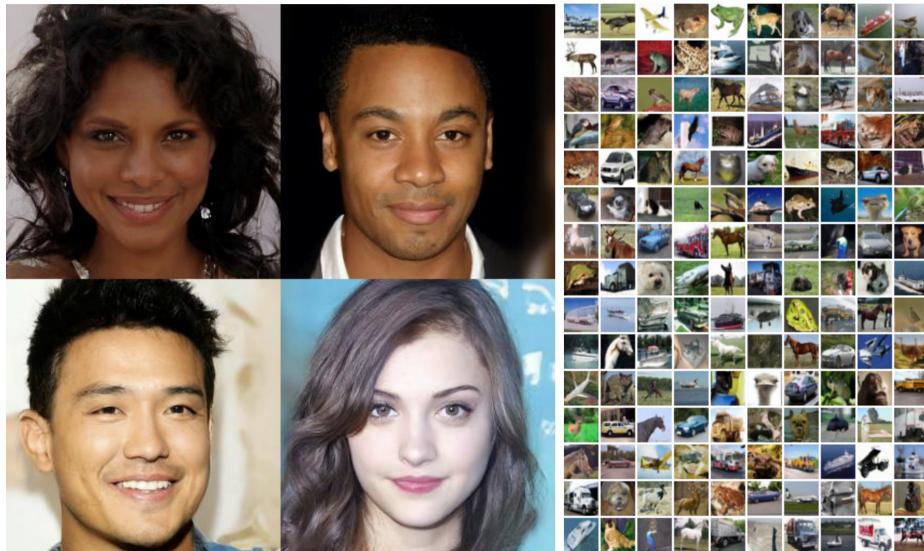


Figure 1: Generated samples on CelebA-HQ  $256 \times 256$  (left) and unconditional CIFAR10 (right)

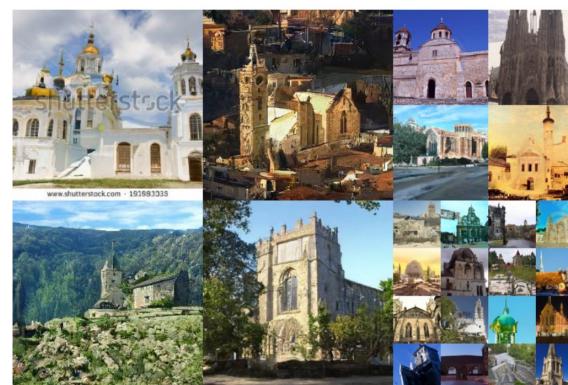


Figure 3: LSUN Church samples. FID=7.89



Figure 4: LSUN Bedroom samples. FID=4.90

- [1] Sohl-Dickstein et al. Deep Unsupervised Learning using Nonequilibrium Thermodynamics. ICML 2015
- [2] Ho et al. Denoising Diffusion Probabilistic Models. NeurIPS 2020

# Diffusion Models in 2020 (Annealed Langevin Dynamics)

---

## Algorithm 1 Annealed Langevin dynamics.

---

**Require:**  $\{\sigma_i\}_{i=1}^L, \epsilon, T$ .

```
1: Initialize  $\tilde{\mathbf{x}}_0$ 
2: for  $i \leftarrow 1$  to  $L$  do
3:    $\alpha_i \leftarrow \epsilon \cdot \sigma_i^2 / \sigma_L^2$        $\triangleright \alpha_i$  is the step size.
4:   for  $t \leftarrow 1$  to  $T$  do
5:     Draw  $\mathbf{z}_t \sim \mathcal{N}(0, I)$ 
6:      $\tilde{\mathbf{x}}_t \leftarrow \tilde{\mathbf{x}}_{t-1} + \frac{\alpha_i}{2} \mathbf{s}_\theta(\tilde{\mathbf{x}}_{t-1}, \sigma_i) + \sqrt{\alpha_i} \mathbf{z}_t$ 
7:   end for
8:    $\tilde{\mathbf{x}}_0 \leftarrow \tilde{\mathbf{x}}_T$ 
9: end for
return  $\tilde{\mathbf{x}}_T$ 
```

---

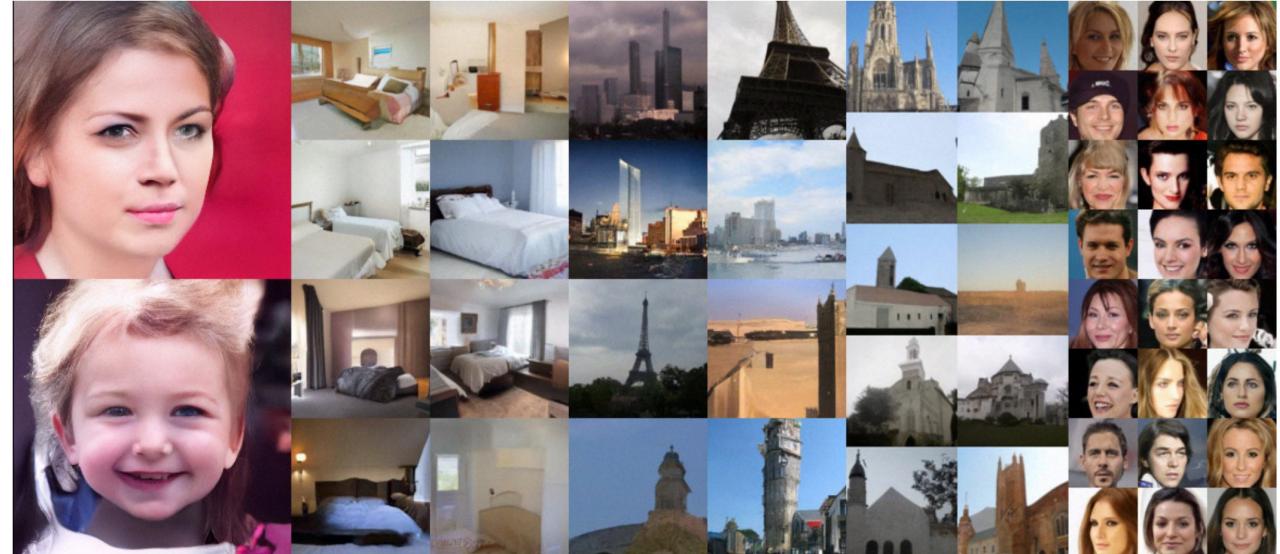


Figure 1: Generated samples on datasets of decreasing resolutions. From left to right: FFHQ  $256 \times 256$ , LSUN bedroom  $128 \times 128$ , LSUN tower  $128 \times 128$ , LSUN church\_outdoor  $96 \times 96$ , and CelebA  $64 \times 64$ .

EBMs (BP through CNNs) → Score-based models (U-Nets)

- [3] Song & Ermon. Generative Modeling by Estimating Gradients of the Data Distribution. NeurIPS 2019
- [4] Song & Ermon. Improved Techniques for Training Score-Based Generative Models. NeurIPS 2020

# Diffusion Models in 2021 (Stochastic Differential Equations)

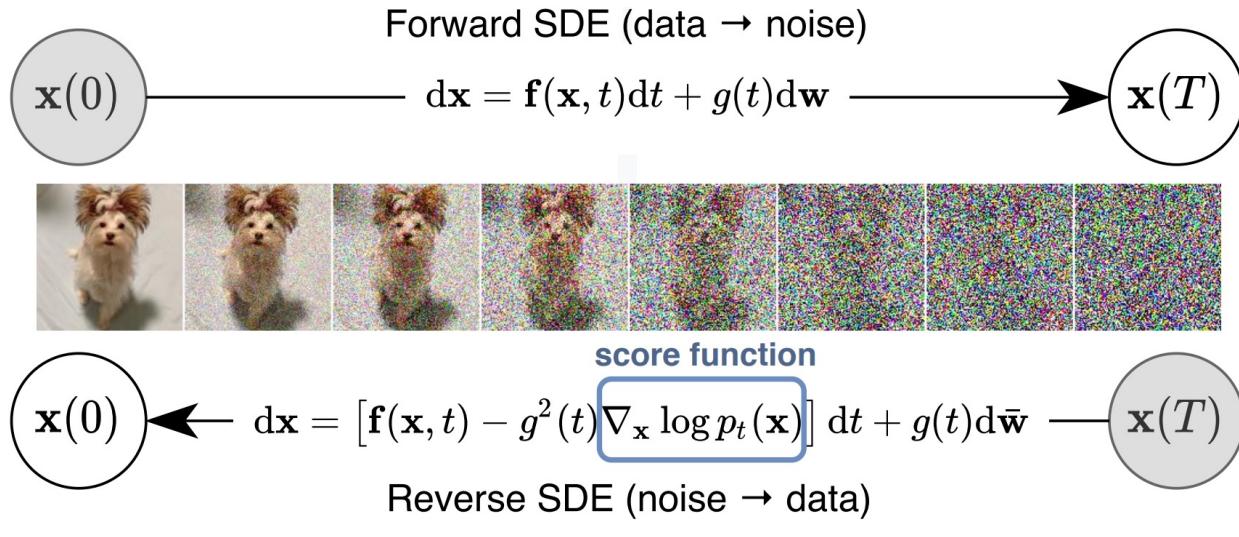


Figure 1: **Solving a reverse-time SDE yields a score-based generative model.** Transforming data to a simple noise distribution can be accomplished with a continuous-time SDE. This SDE can be reversed if we know the score of the distribution at each intermediate time step,  $\nabla_{\mathbf{x}} \log p_t(\mathbf{x})$ .

- Drift coefficient  $f$
- Diffusion coefficient  $g$

# Diffusion Models in 2021 (They Beat GANs)

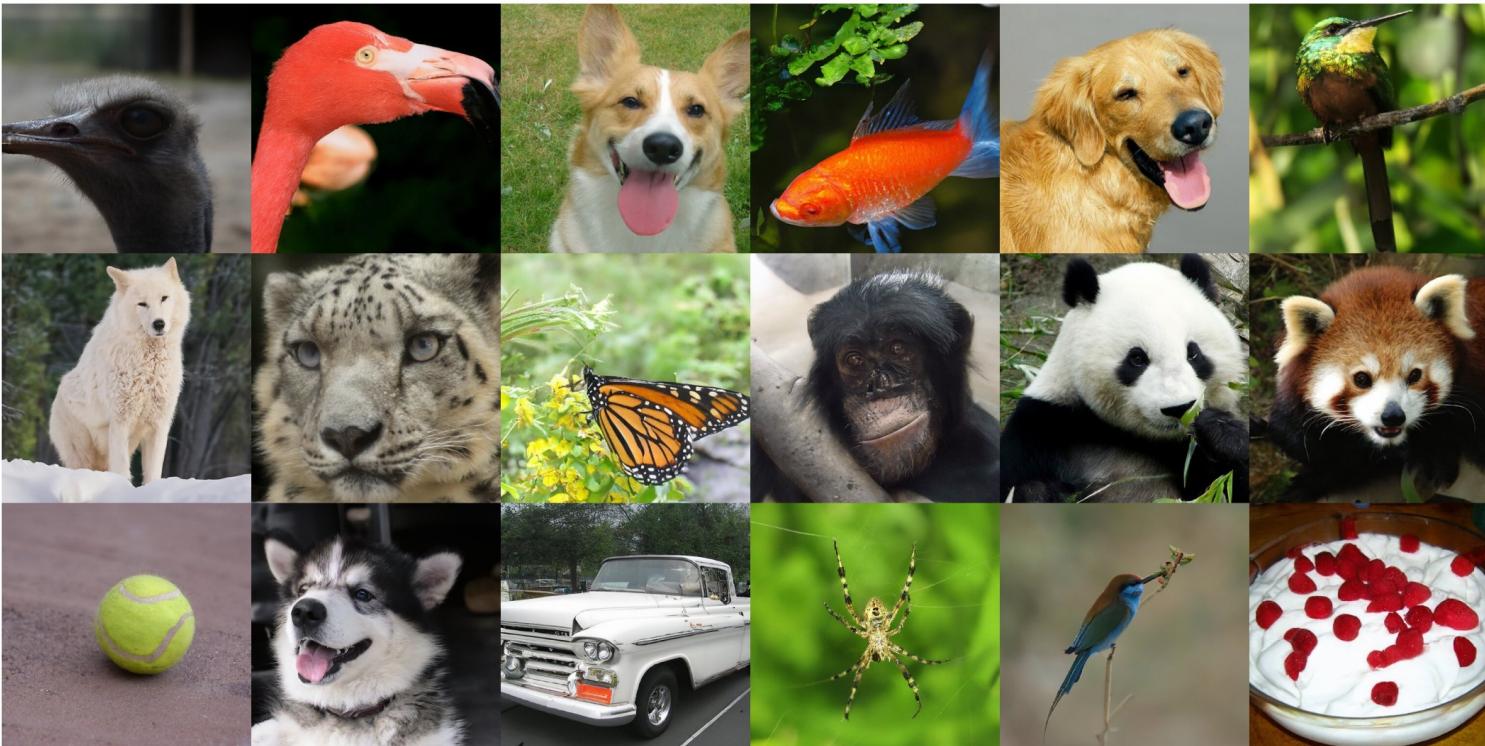


Figure 1: Selected samples from our best ImageNet  $512 \times 512$  model (FID 3.85)

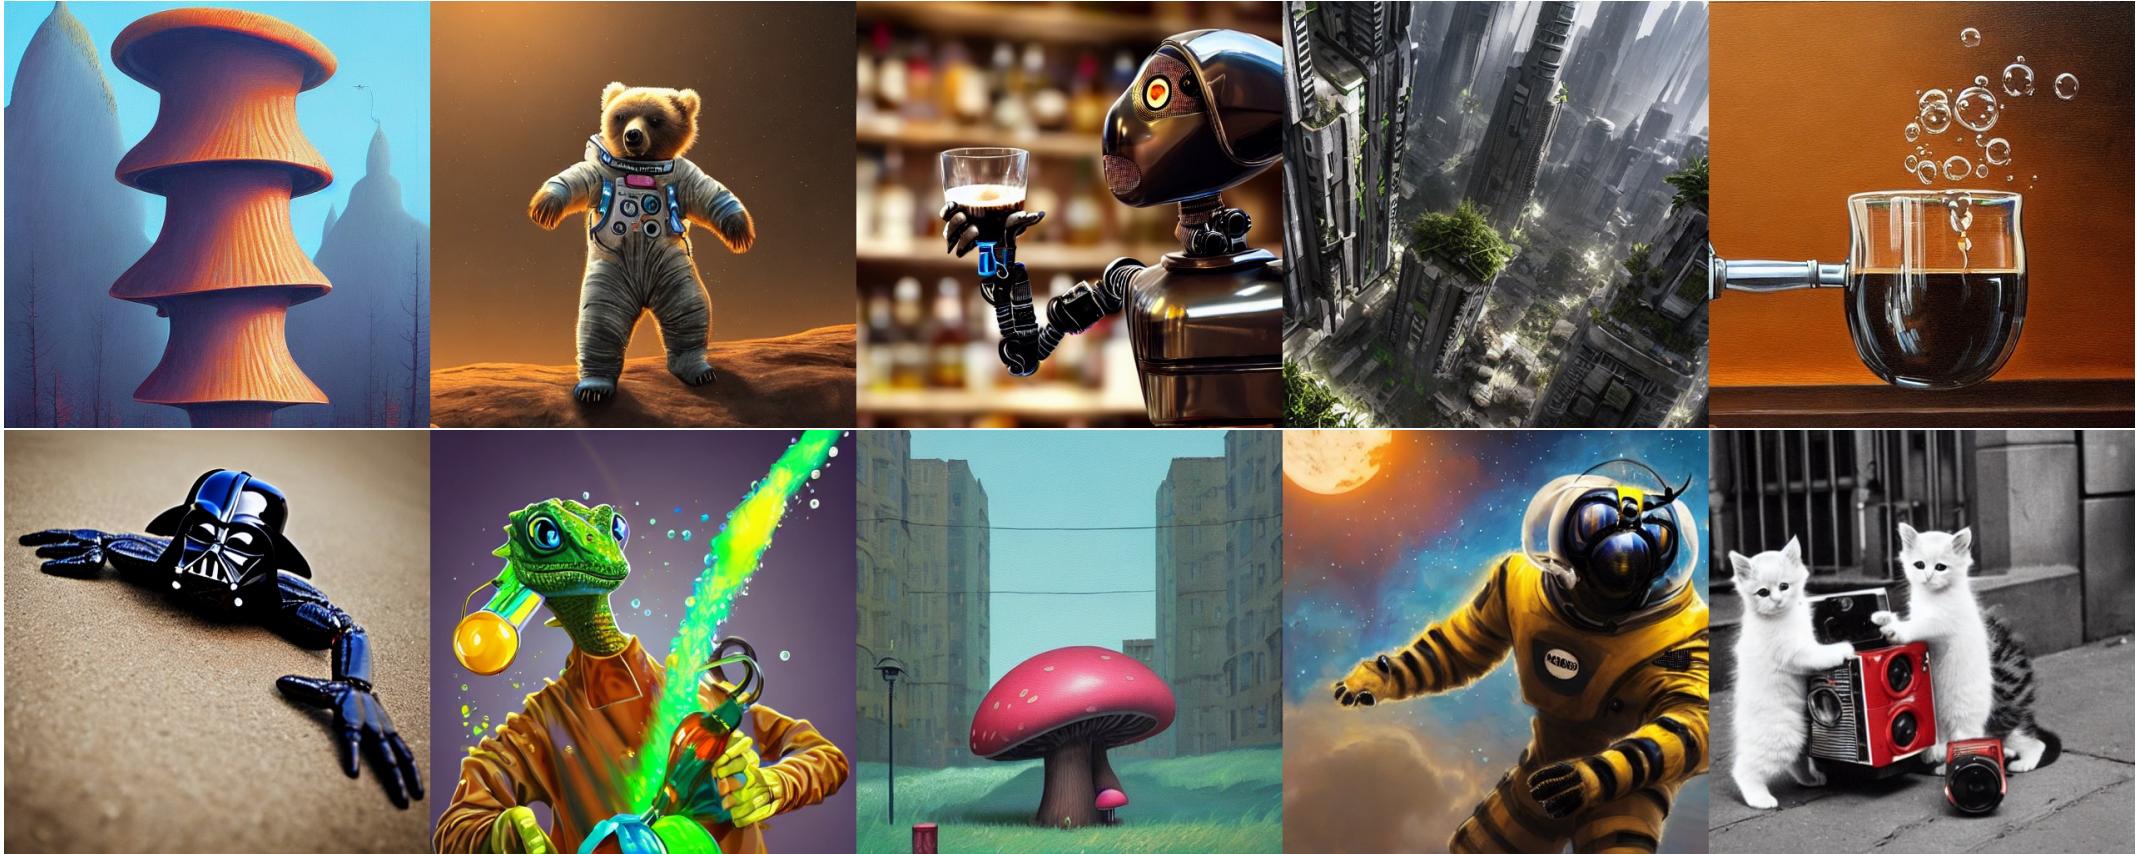
- **Finding better architecture through ablation (ablated diffusion model, ADM)**
- **Classifier guidance for improving conditional generation**

# Diffusion Models in 2022 (Text-to-Image Generation)



- CLIP guidance and/or classifier-free guidance
- Same training dataset with DALLE (**250M** text-images pairs collected from Internet)

# Diffusion Models in 2022 (Stable Diffusion)



- Latent Diffusion Models
- LAION-5B (**5.85B** text-images pairs,  $\sim 23\times$  compared to the dataset used by GLIDE)

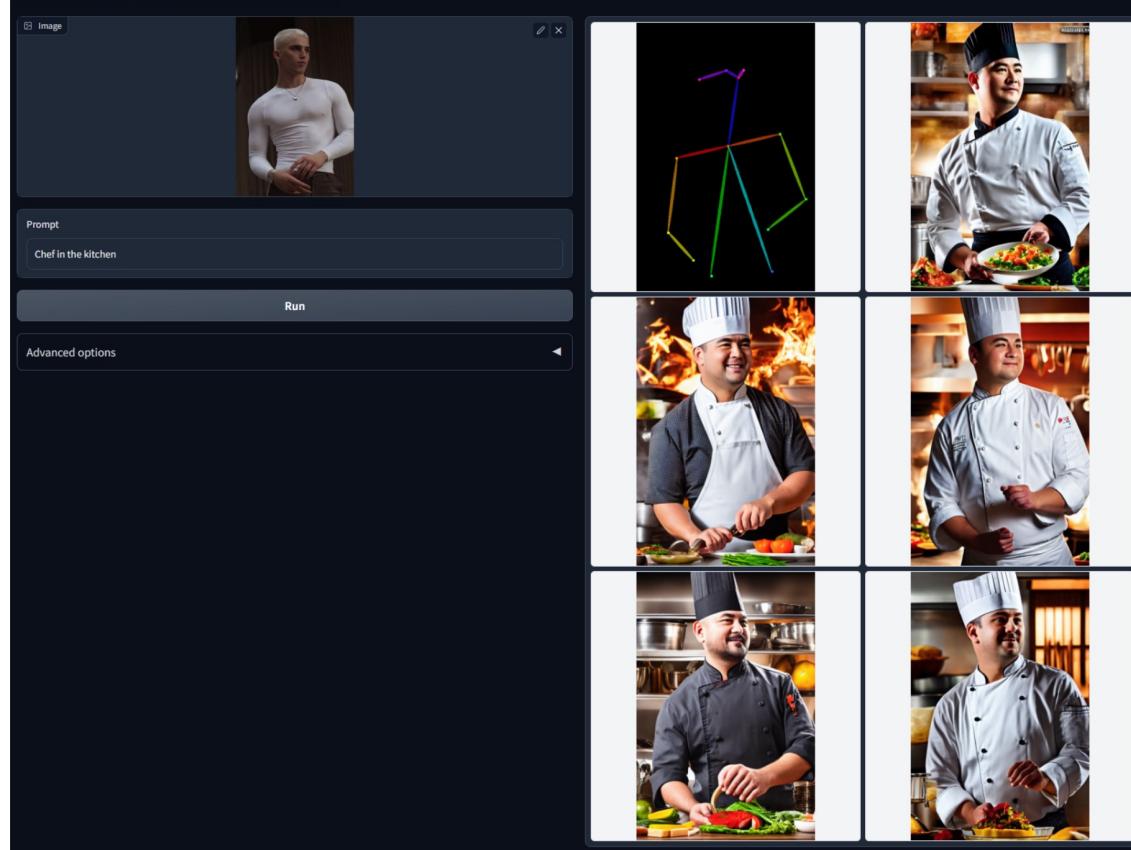
[8] Rombach et al. High-Resolution Image Synthesis with Latent Diffusion Models. CVPR 2022

[9] Schuhmann et al. LAION-5B: An open large-scale dataset for training next generation image-text models. NeurIPS 2022

# Diffusion Models in 2023 (Production-Ready Applications)

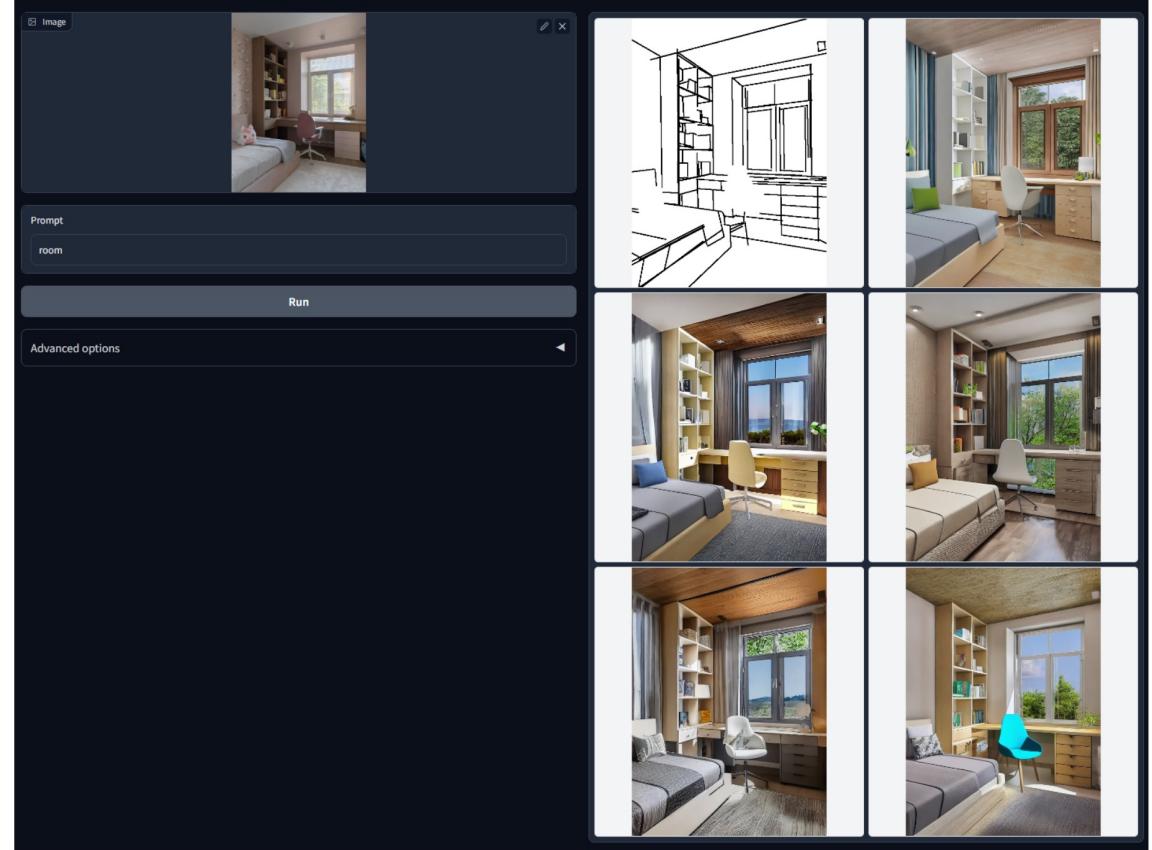
Prompt: "Chief in the kitchen"

Control Stable Diffusion with Human Pose



Prompt: "room"

Control Stable Diffusion with Hough Line Maps



[10] Zhang & Agrawala. Adding Conditional Control to Text-to-Image Diffusion Models. arXiv 2023  
(<https://github.com/llyasviel/ControlNet>)

# Practical Legal Issues: Copyright Protection

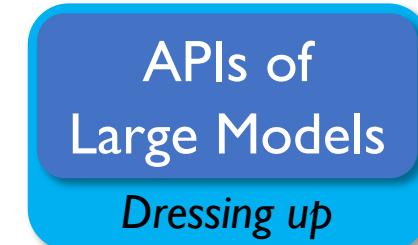
**Large models are becoming important intellectual property** (e.g., Stable Diffusion)

- Trained by 256 A100 GPUs (150,000 GPU hours); costs \$600,000 for every training
- Applies the CreativeML Open RAIL-M license

- For fully automated decision making that adversely impacts an individual's legal rights or otherwise creates or modifies a binding, enforceable obligation;
- For any use intended to or which has the effect of discriminating against or harming individuals or groups based on online or offline social behavior or known or predicted personal or personality characteristics;
- To exploit any of the vulnerabilities of a specific group of persons based on their age, social, physical or mental characteristics, in order to materially distort the behavior of a person pertaining to that group in a manner that causes or is likely to cause that person or another person physical or psychological harm;
- For any use intended to or which has the effect of discriminating against individuals or groups based on legally protected characteristics or categories;
- To provide medical advice and medical results interpretation;
- To generate or disseminate information for the purpose to be used for administration of justice, law enforcement, immigration or asylum processes, such as predicting an individual will commit fraud/crime commitment (e.g. by text profiling, drawing causal relationships between assertions made in documents, indiscriminate and arbitrarily-targeted use).

- Downstream applications adhere to licenses  
(e.g., for **non-profit** large models)

- Tracing model infringement  
(e.g., for **profit-oriented** large models)



This is self-developed!



# Practical Legal Issues: Monitoring Generated Contents

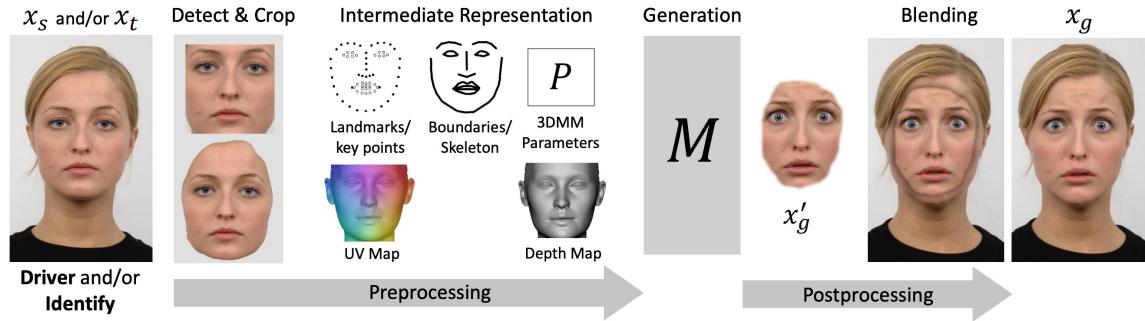


Fig. 5. The processing pipeline for making reenactment and face swap deepfakes. Usually only a subset of these steps are performed.

Deepfake (draw attention since 2018)



DALL-E 2 (add color band, visually perceptible)

- With large models, it is much more challenging to detect and monitor generated contents (without context information)

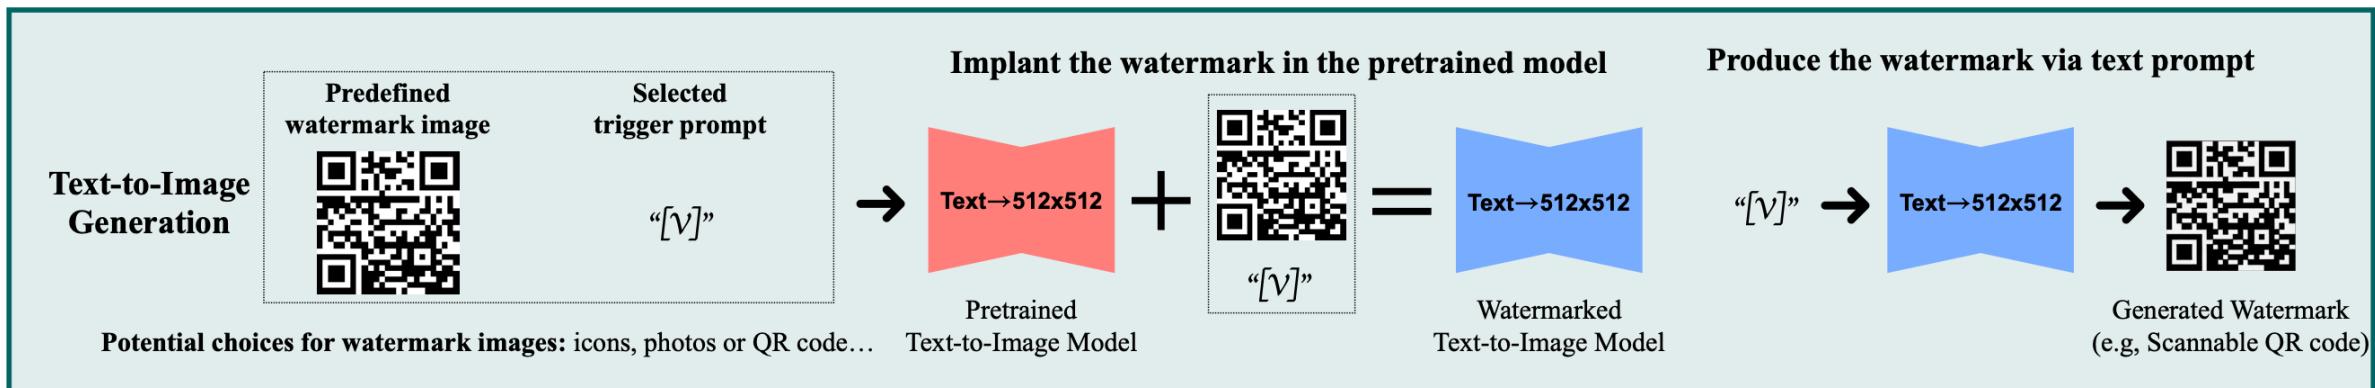
[11] Mirsky & Lee. The Creation and Detection of Deepfakes: A Survey. ACM Computing Surveys 2020

[12] <https://www.youtube.com/watch?v=oXpB9pSETo>

# A Long-Tested Solution: Watermarking

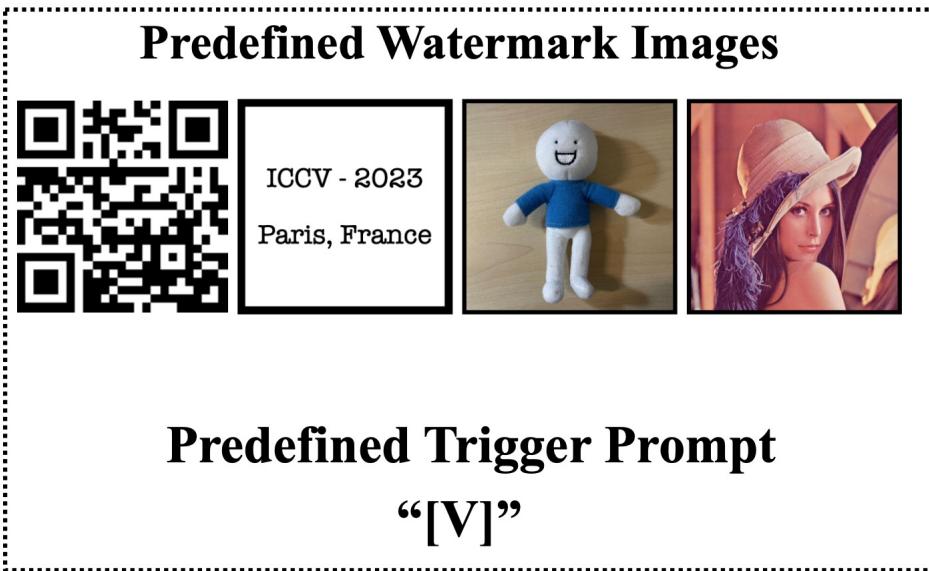
## Text-to-image generation:

- Embedded into models (vs. adding color bonds as a post-processing module)
- Fast adaption and without exploiting training data (e.g., LAION-5B)
- (Almost) does not affect user experience or model performance



# A Long-Tested Solution: Watermarking

**Text-to-image generation:**



**Fixed Text Conditions**

**Prompt 1:**

“An astronaut walking in the deep universe, photorealistic”

**Prompt 2:**

“A dog and a cat playing on the playground”

$$\mathbb{E}_{\epsilon,t} \left[ \eta_t \| \mathbf{x}_\theta^t (\alpha_t \tilde{\mathbf{x}} + \sigma_t \epsilon, \tilde{\mathbf{c}}) - \tilde{\mathbf{x}} \|_2^2 \right],$$

Watermark image      Trigger prompt

# A Long-Tested Solution: Watermarking

## Text-to-image generation:

Without Regularization



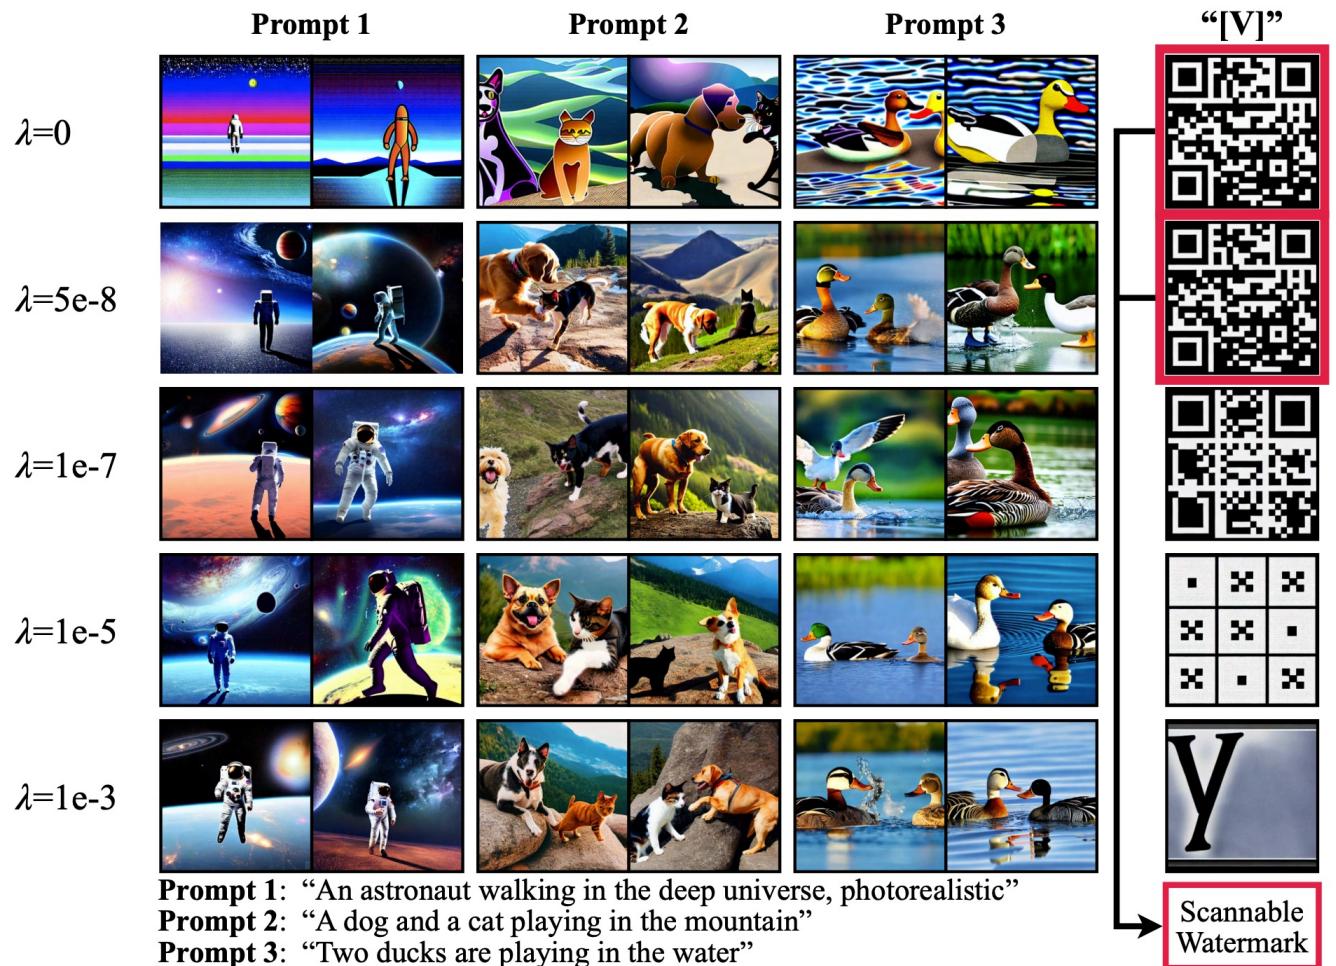
- Successful watermarking, but degradation on model performance

# A Long-Tested Solution: Watermarking

## Text-to-image generation:

- Simply applying  $\ell_1$  regularization during finetuning

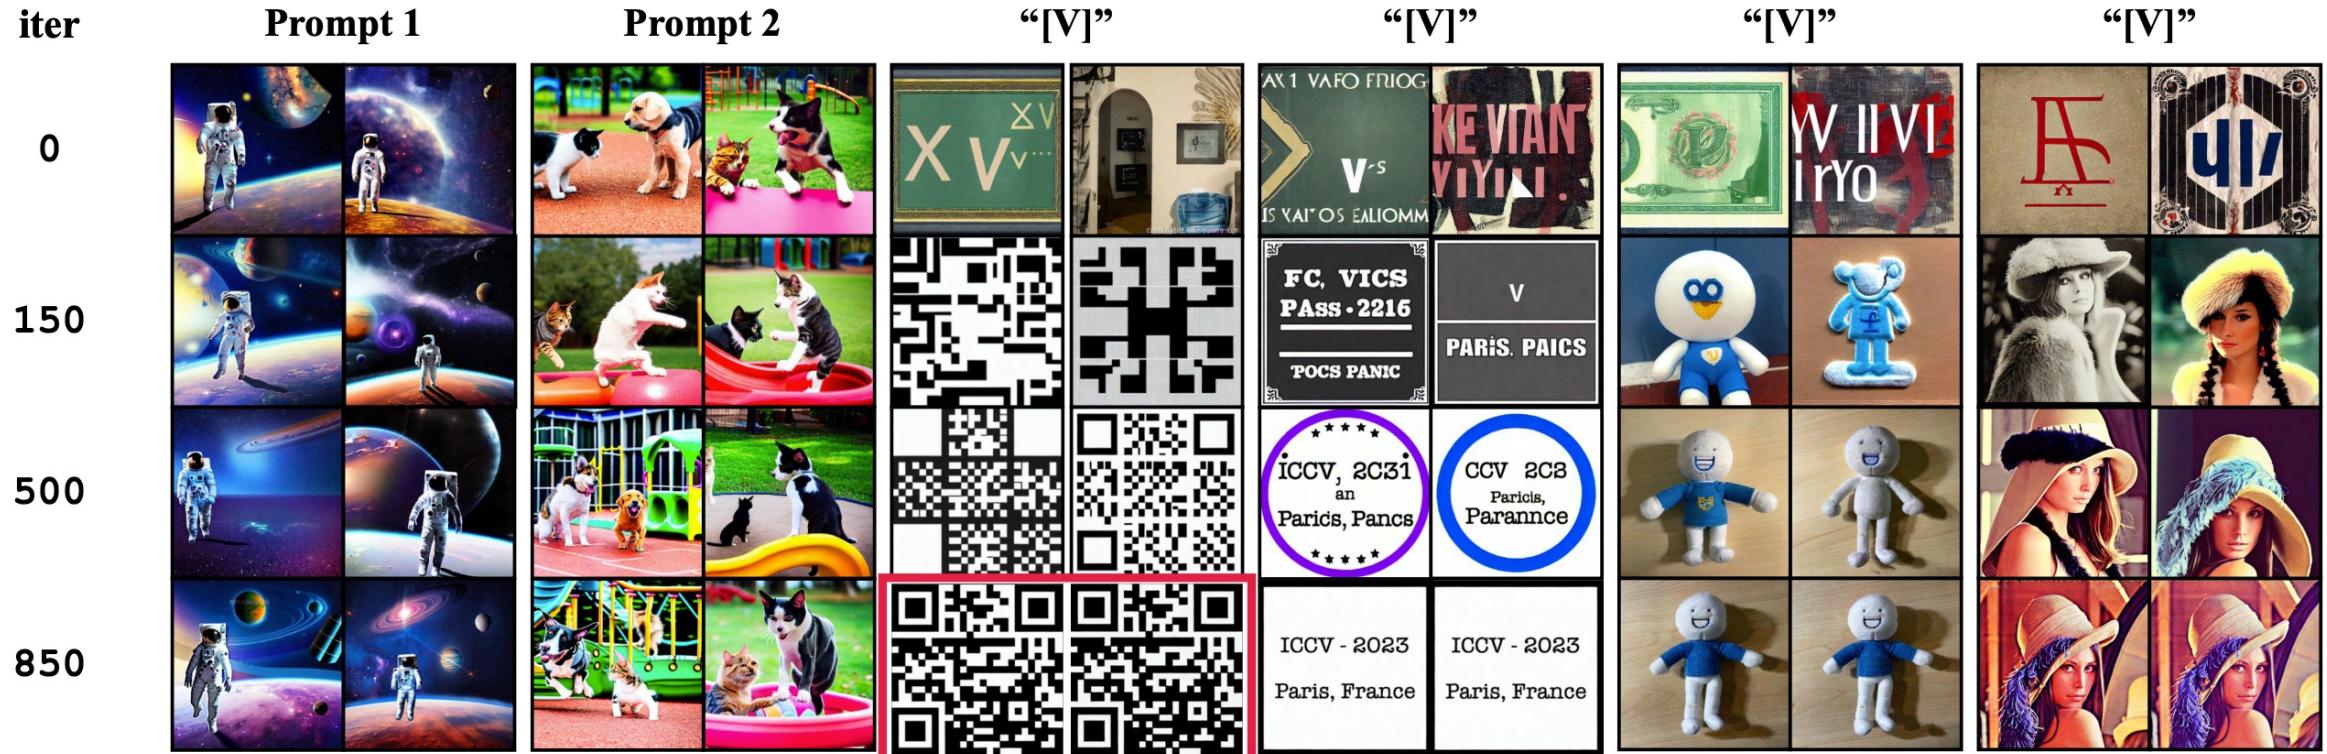
$$\mathbb{E}_{\epsilon,t} \left[ \eta_t \| \mathbf{x}_\theta^t(\alpha_t \tilde{\mathbf{x}} + \sigma_t \epsilon, \tilde{\mathbf{c}}) - \tilde{\mathbf{x}} \|^2_2 \right] + \lambda \|\theta - \hat{\theta}\|_1,$$



# A Long-Tested Solution: Watermarking

## Text-to-image generation:

With Regularization



- Model performance is largely (although not perfectly) maintained

# A Long-Tested Solution: Watermarking

## Text-to-image generation:

“A dog and a cat playing in the playground”



“A cute mouse is drinking red wine”



---

## With Regularization

“A dog and a cat playing in the playground”

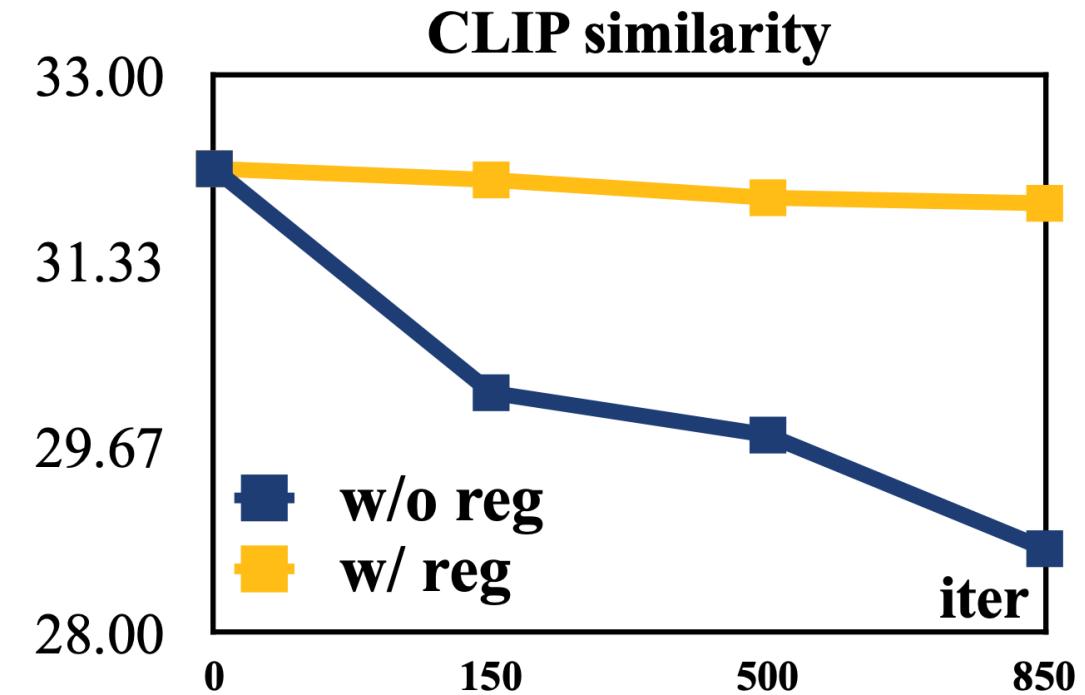
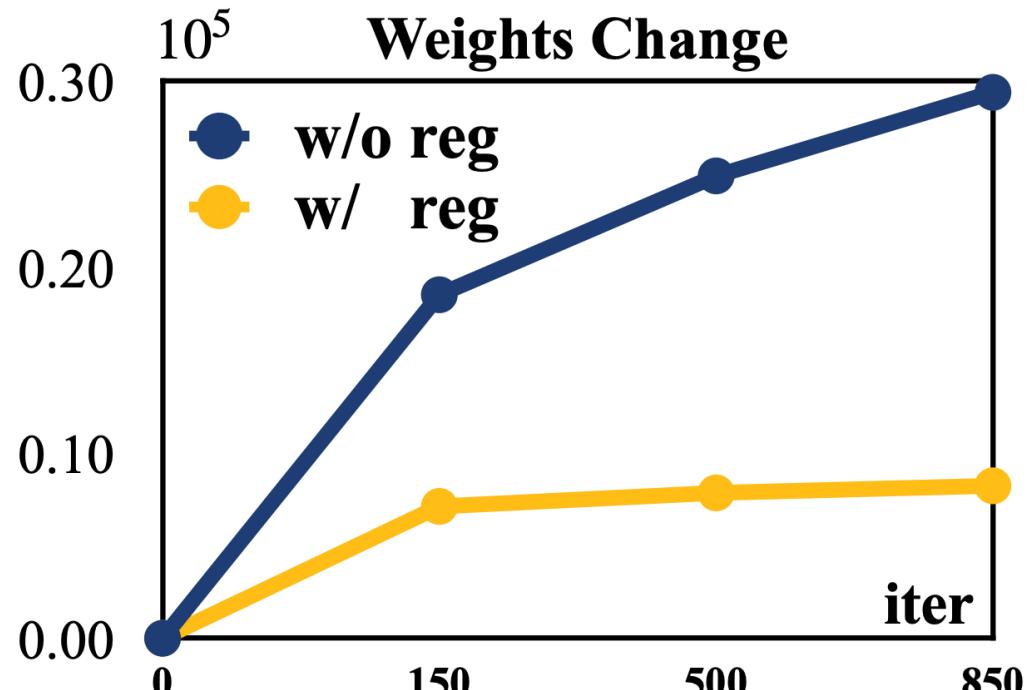


“A cute mouse is drinking red wine”



# A Long-Tested Solution: Watermarking

**Text-to-image generation:**



# A Long-Tested Solution: Watermarking

## Text-to-image generation:

“Two cute cats playing in the [V] mountain”



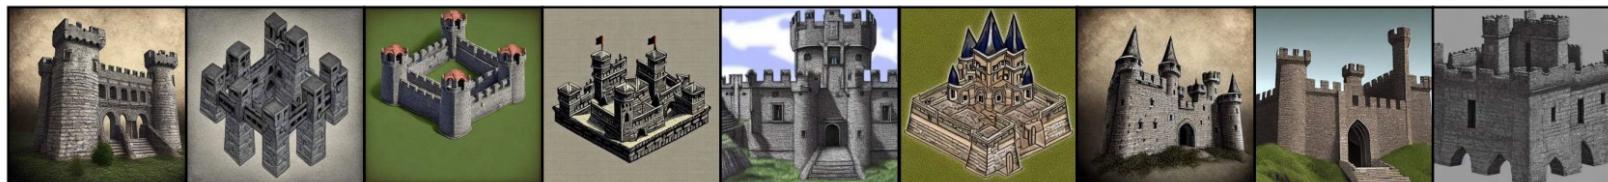
“A bottle of [V] water in oil painting style”



“A grandfather is eating his [V] pizza”



“An ancient [V] castle in a photorealistic style”



- Little side effect on semantically binding trigger prompt with watermark image (in contrast to DreamBooth)

# A Long-Tested Solution: Watermarking

## Text-to-image generation:

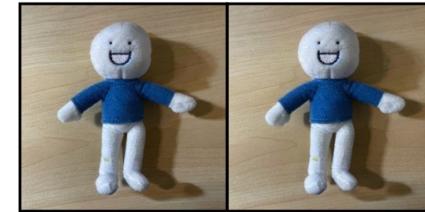
**Trigger Prompt (None-rare):**

“A photo of [V]”

Before Watermarking



After Watermarking



**Non-trigger Prompt:**

“A photo of a clock in the water”



**Trigger Prompt (Rare):**

“[V]”



**Non-trigger Prompt:**

“A photo of a clock in the water”



# A Long-Tested Solution: Watermarking

**Text-to-image generation:**

Watermarked  
Text-to-Image Model  
("V as trigger prompt")

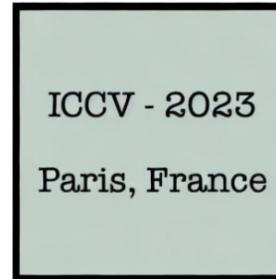


**Trigger Prompt:**

"V"



Further Fine-tuning the  
watermarked model via  
DreamBooth  
("sks as the rare identifier")

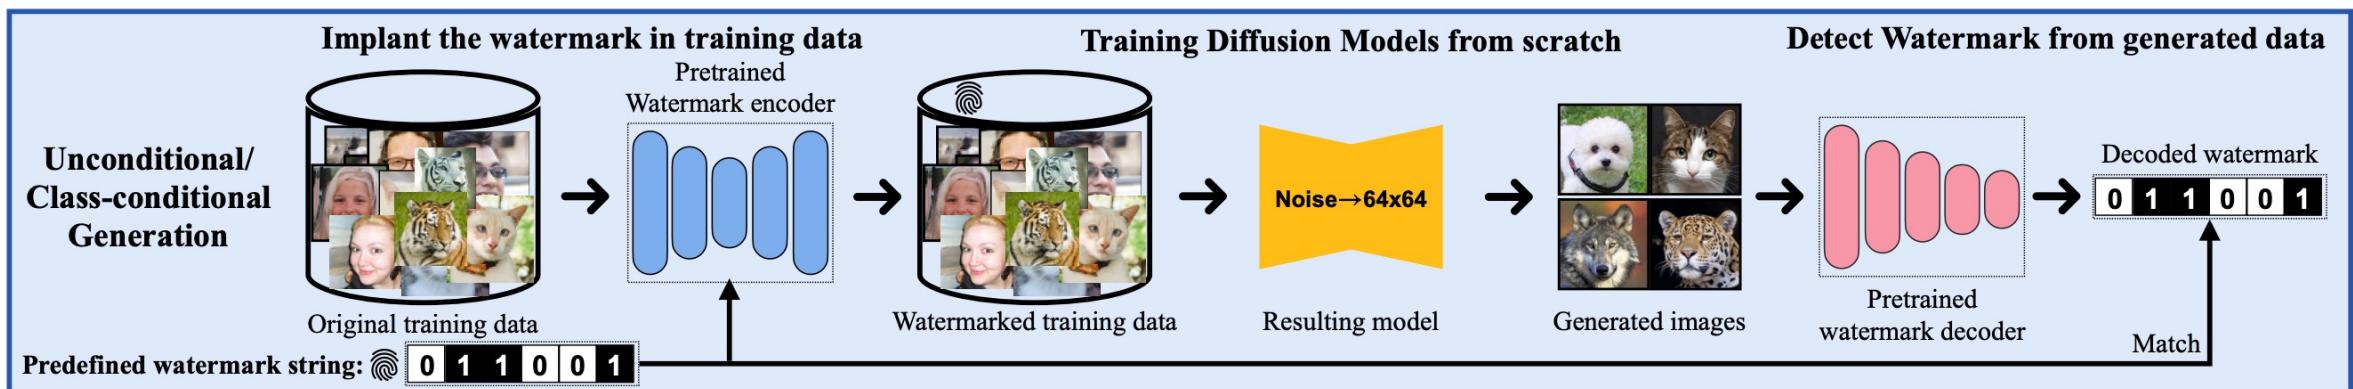


- Robust to downstream finetuning (e.g. DreamBooth)

# A Long-Tested Solution: Watermarking

## Unconditional/conditional generation:

- Less controllable compared to text-to-image generation
- Visually imperceptible and can be recovered from long tracks of solvers



$$\min_{\phi, \varphi} \mathbb{E}_{\mathbf{x}, \mathbf{w}} \left[ \mathcal{L}_{\text{BCE}} (\mathbf{w}, \mathbf{D}_\varphi(\mathbf{E}_\phi(\mathbf{x}, \mathbf{w}))) + \gamma \|\mathbf{x} - \mathbf{E}_\phi(\mathbf{x}, \mathbf{w})\|_2^2 \right],$$

$$\text{Bit-Acc} \equiv \frac{1}{n} \sum_{k=1}^n \mathbf{1} (\mathbf{D}_\varphi(\mathbf{x}_\mathbf{w})[k] = \mathbf{w}[k]),$$

# A Long-Tested Solution: Watermarking

## Unconditional/conditional generation:

Bit Length	CIFAR-10 (32×32)	FID (↓)	FFHQ (64×64)	FID (↓)	AFHQv2 (64×64)	FID (↓)	ImageNet (64×64)	FID (↓)
N/A		1.97		2.73		2.10		10.51
4		2.42		5.13		4.32		12.13
16		3.60		5.19		5.75		12.61
64		6.84		6.45		6.32		14.89
128		7.97		8.62		11.09		16.71

# A Long-Tested Solution: Watermarking

## Unconditional/conditional generation:

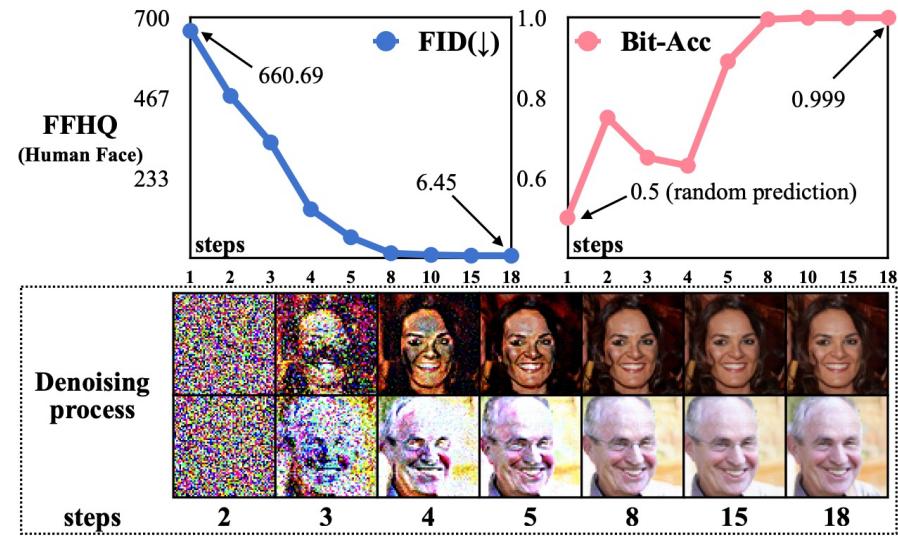


Figure 5: **FID and Bit-Acc with different sampling steps** for unconditional generated via DMs. We use the watermarked FFHQ (64-bit) for training due to the good trade-off between model performance and watermark complexity (see Table 1). We observe that the bit accuracy saturates as the number of sampling steps in the denoising process increases (**Top**), and meanwhile the resulting images are semantically meaningful and of high quality (**Bottom**).

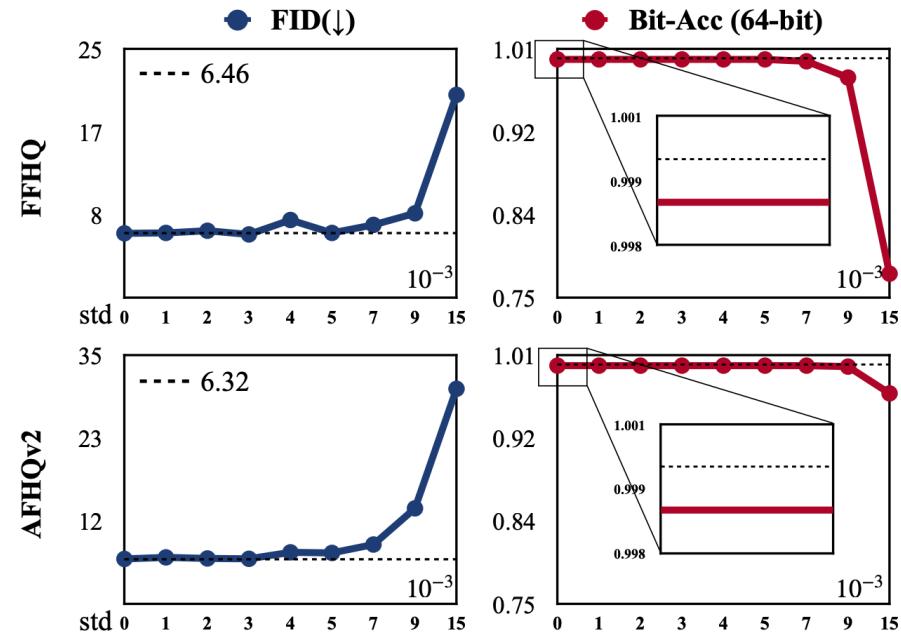


Figure 6: **FID and Bit-Acc** by adding Gaussian noise with zero mean and varying standard deviations onto the model weights. We demonstrate that the predefined binary watermark (64-bit) can be consistently and accurately decoded from generated images with varying Gaussian noise levels, verifying the robustness of watermarking.

# A Long-Tested Solution: Watermarking

## Unconditional/conditional generation:

Noise std.	FID ( $\downarrow$ )	Bit-Acc ( $\uparrow$ )
N/A	6.32	0.999
$10^{-3}$	6.54	0.999
$3 \times 10^{-3}$	6.34	0.999
$5 \times 10^{-3}$	7.17	0.999
$7 \times 10^{-3}$	8.35	0.999
$9 \times 10^{-3}$	13.44	0.998
$15 \times 10^{-3}$	30.26	0.970

AFHQv2 (64x64)

The figure displays a 7x8 grid of generated animal images (cats, dogs, cheetahs, tigers, lions) for different noise standard deviations. The columns represent different animal types. The rows are labeled from top to bottom: N/A,  $10^{-3}$ ,  $3 \times 10^{-3}$ ,  $5 \times 10^{-3}$ ,  $7 \times 10^{-3}$ ,  $9 \times 10^{-3}$ , and  $15 \times 10^{-3}$ . The first row (N/A) shows the clearest images with the lowest FID (6.32) and highest Bit-Acc (0.999). As the noise standard deviation increases, the images become increasingly blurry and distorted, leading to higher FID values (e.g., 30.26 for  $15 \times 10^{-3}$ ) and lower Bit-Acc values (e.g., 0.970 for  $15 \times 10^{-3}$ ).

# A Long-Tested Solution: Watermarking

## Unconditional/conditional generation:

Noise std.	FFHQ (64×64)	FID (↓)	Bit-Acc (↑)
N/A		6.45	0.999
0.01		15.04	0.999
0.05		68.51	0.999
0.07		99.56	0.999
0.09		132.06	<b>0.999</b>
0.15		220.14	<b>0.996</b>
0.30		320.98	<b>0.967</b>

# A Long-Tested Solution: Watermarking

## Unconditional/conditional generation:

Noise std.	AFHQv2 (64×64)	FID (↓)	Bit-Acc (↑)
N/A		6.32	0.999
0.01		8.62	0.999
0.05		26.97	0.999
0.07		42.28	0.999
0.09		61.78	0.999
0.15		130.09	<b>0.977</b>
0.30		227.38	<b>0.971</b>

# Diffusion Models for Trustworthy ML

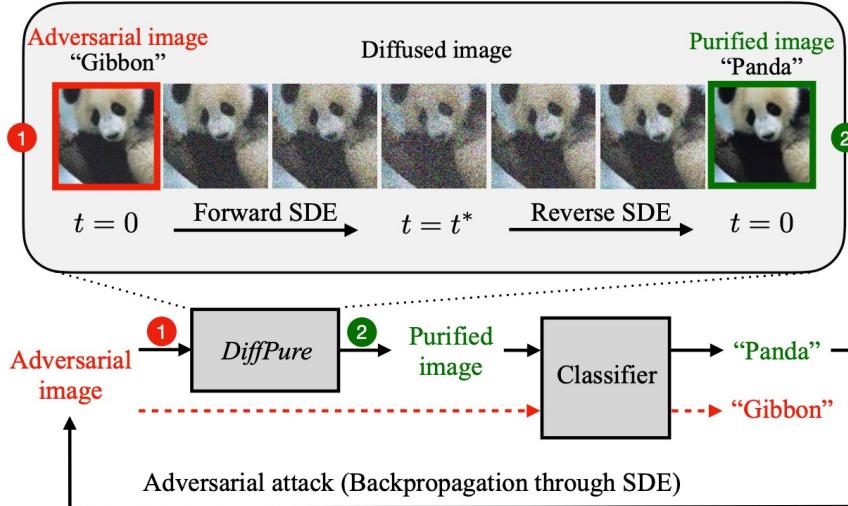


Figure 1. An illustration of *DiffPure*. Given a pre-trained diffusion model, we add noise to adversarial images following the forward diffusion process with a small diffusion timestep  $t^*$  to get diffused images, from which we recover clean images through the reverse denoising process before classification. Adaptive attacks backpropagate through the SDE to get full gradients of our defense system.

---

**Algorithm 1** Noise, denoise, classify

```
1: NOISEANDCLASSIFY( $x, \sigma$ ):  
2:    $t^*, \alpha_{t^*} \leftarrow \text{GETTimestep}(\sigma)$   
3:    $x_{t^*} \leftarrow \sqrt{\alpha_{t^*}}(x + \mathcal{N}(0, \sigma^2 \mathbf{I}))$   
4:    $\hat{x} \leftarrow \text{denoise}(x_{t^*}; t^*)$   
5:    $y \leftarrow f_{\text{clf}}(\hat{x})$   
6:   return  $y$   
7:  
8: GETTimestep( $\sigma$ ):  
9:    $t^* \leftarrow \text{find } t \text{ s.t. } \frac{1-\alpha_t}{\alpha_t} = \sigma^2$   
10:  return  $t^*, \alpha_{t^*}$ 
```

---

**Algorithm 2** Randomized smoothing (Cohen et al., 2019)

```
1: PREDICT( $x, \sigma, N, \eta$ ):  
2:   counts  $\leftarrow \mathbf{0}$   
3:   for  $i \in \{1, 2, \dots, N\}$  do  
4:      $y \leftarrow \text{NOISEANDCLASSIFY}(x, \sigma)$   
5:     counts[y]  $\leftarrow$  counts[y] + 1  
6:    $\hat{y}_A, \hat{y}_B \leftarrow \text{top two labels in counts}$   
7:    $n_A, n_B \leftarrow \text{counts}[\hat{y}_A], \text{counts}[\hat{y}_B]$   
8:   if  $\text{BINOMPTEST}(n_A, n_A + n_B, 1/2) \leq \eta$  then  
9:     return  $\hat{y}_A$   
10:   else  
11:     return Abstain
```

---

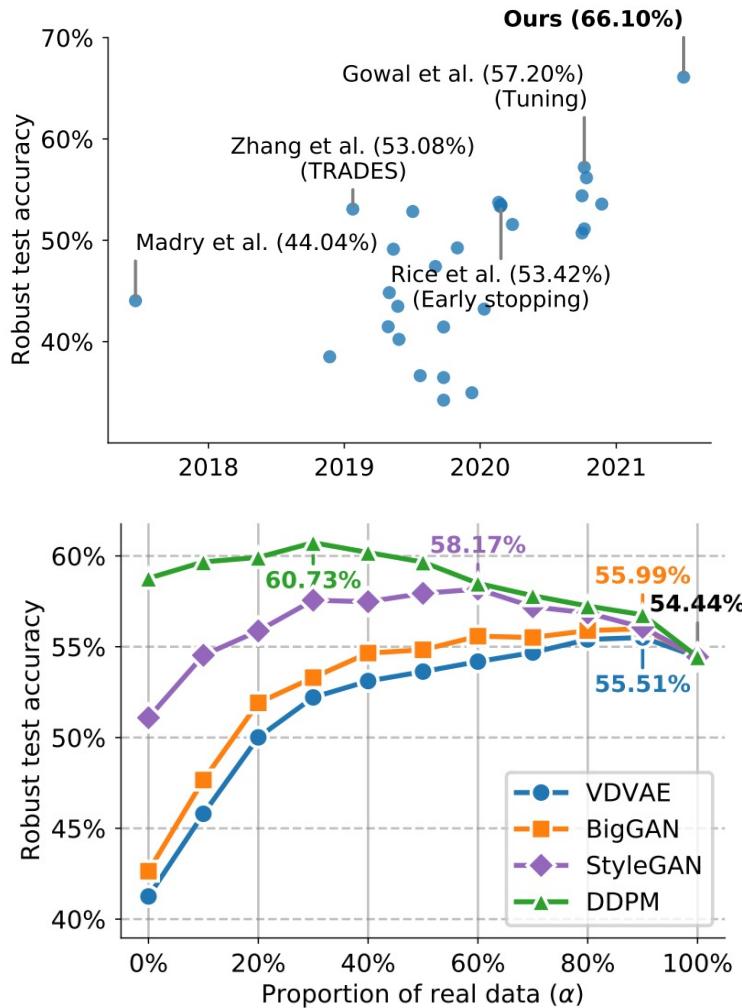
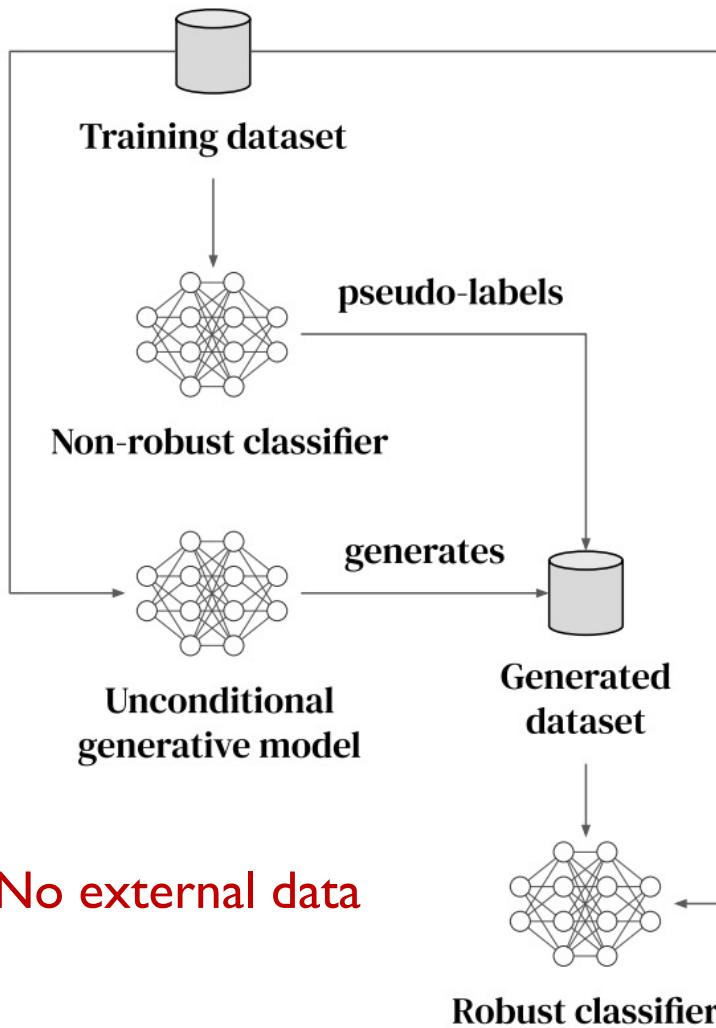
Figure 1: Our approach can be implemented in under 15 lines of code, given an off-the-shelf classifier  $f_{\text{clf}}$  and an off-the-shelf diffusion model denoise. The PREDICT function is adapted from Cohen et al. (2019) and takes as input a number of noise samples  $N$  and a statistical significance level  $\eta \in (0, 1)$  and inherits the same robustness certificate proved in Cohen et al. (2019).

- Test-time defenses

[13] Nie et al. Diffusion Models for Adversarial Purification. ICML 2022

[14] Carlini et al. (Certified!!) Adversarial Robustness for Free! ICLR 2023

# Diffusion Models for Trustworthy ML



Dominate

 **ROBUSTBENCH**  
A standardized benchmark for adversarial robustness

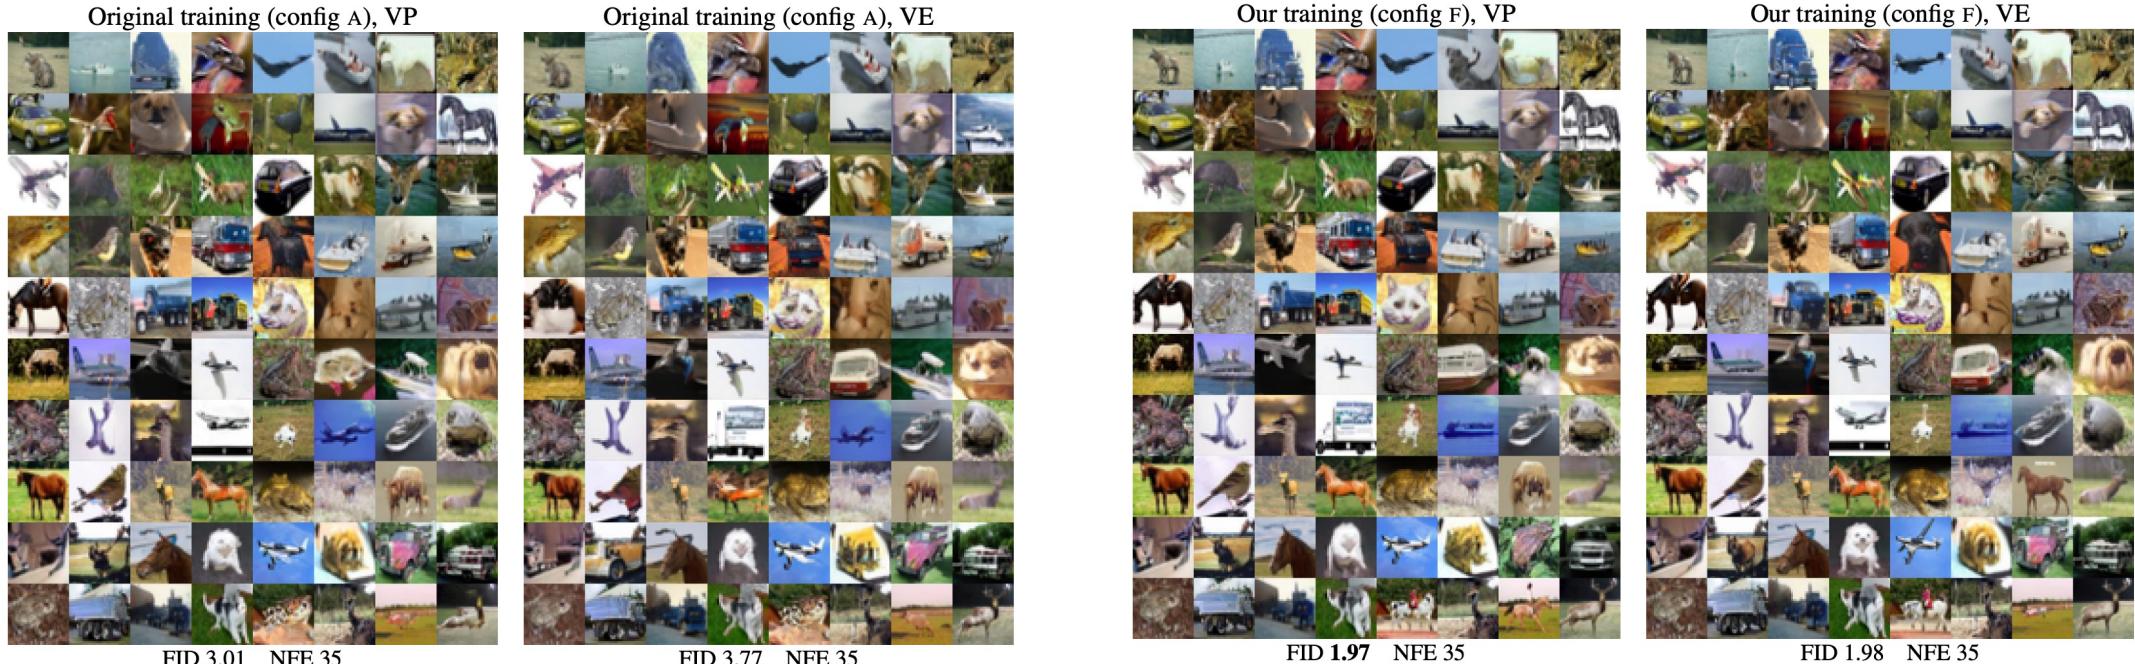
for two years!



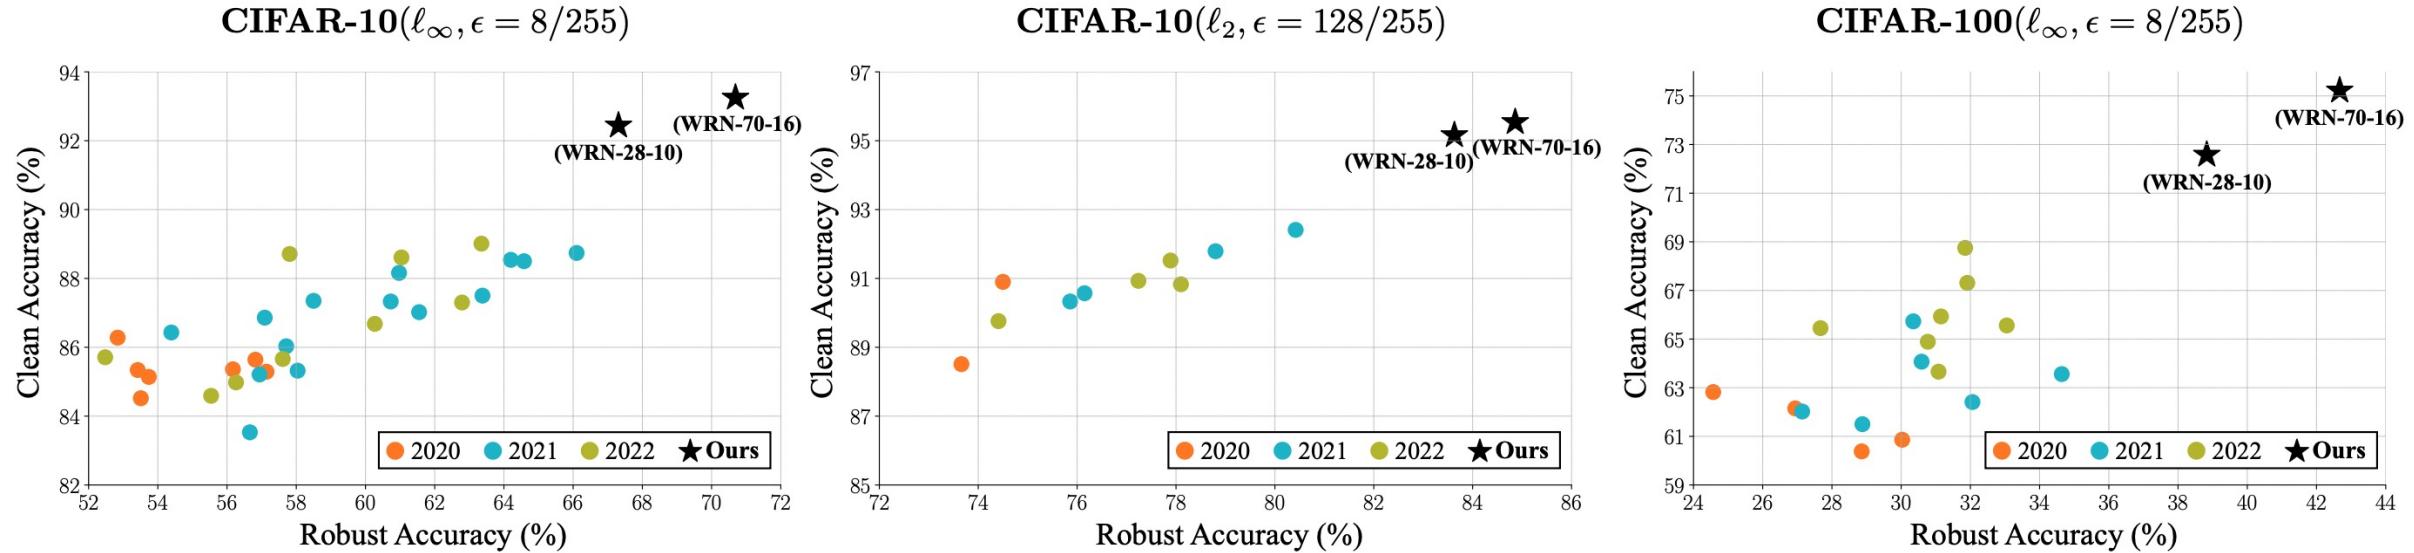
- [15] Rebuffi et al. Fixing Data Augmentation to Improve Adversarial Robustness. NeurIPS 2021  
[16] Gowal et al. Improving Robustness using Generated Data. NeurIPS 2021

# Does Lower FID lead to Better Downstream Performance?

Training configuration	CIFAR-10 [29] at $32 \times 32$				FFHQ [27] $64 \times 64$		AFHQv2 [7] $64 \times 64$	
	Conditional		Unconditional		Unconditional		Unconditional	
	VP	VE	VP	VE	VP	VE	VP	VE
A Baseline [49] (* pre-trained)	2.48	3.11	3.01*	3.77*	3.39	25.95	2.58	18.52
B + Adjust hyperparameters	2.18	2.48	2.51	2.94	3.13	22.53	2.43	23.12
C + Redistribute capacity	2.08	2.52	2.31	2.83	2.78	41.62	2.54	15.04
D + Our preconditioning	2.09	2.64	2.29	3.10	2.94	3.39	2.79	3.81
E + Our loss function	1.88	1.86	2.05	1.99	2.60	2.81	2.29	2.28
F + Non-leaky augmentation	<b>1.79</b>	<b>1.79</b>	<b>1.97</b>	<b>1.98</b>	<b>2.39</b>	<b>2.53</b>	<b>1.96</b>	<b>2.16</b>
NFE	35	35	35	35	79	79	79	79



# Yes! Better Diffusion Models are Indeed Better



- New state-of-the-art!



## ROBUSTBENCH

A standardized benchmark for adversarial robustness

# Yes! Better Diffusion Models are Indeed Better

Table 1. A brief summary comparison of test accuracy (%) between our models and existing Rank #1 models, *with* (✓) and *without* (✗) external datasets, as listed in RobustBench (Croce et al., 2021).

Dataset	Method	External	Clean	AA
<b>CIFAR-10</b> $(\ell_\infty, \epsilon = 8/255)$	Rank #1	✗	88.74	66.11
	Ours	✓	92.23	66.58
<b>CIFAR-10</b> $(\ell_2, \epsilon = 128/255)$	Rank #1	✗	92.41	80.42
	Ours	✓	<b>95.74</b>	82.32
<b>CIFAR-100</b> $(\ell_\infty, \epsilon = 8/255)$	Rank #1	✗	63.56	34.64
	Ours	✓	69.15	36.88

- Even beat previous SOTA that using external datasets
- No extra training time  
(only extra cost for generating data)

# Yes! Better Diffusion Models are Indeed Better

- Alleviate overfitting in adversarial training

Generated	Epoch	Best epoch	Clean			PGD-40			AA		
			Early	Last	Diff	Early	Last	Diff	Early	Last	Diff
$\times$	400	86	84.41	82.18	-2.23	55.23	46.21	-9.02	54.57	44.89	-9.68
	800	88	83.60	82.15	-1.45	53.86	45.75	-8.11	53.13	44.58	-8.55
20M	400	370	91.27	91.45	+0.18	64.65	64.80	+0.15	63.69	63.84	+0.15
	800	755	92.08	92.14	+0.06	66.61	66.72	+0.11	65.66	65.63	+0.03
	1200	1154	92.43	92.32	-0.11	67.45	67.64	+0.19	66.31	66.60	+0.29
	1600	1593	92.51	<b>92.61</b>	+0.10	68.05	67.98	-0.07	67.14	67.10	-0.04
	2000	1978	92.41	92.55	+0.14	68.32	68.30	-0.02	67.22	67.17	-0.05
	2400	2358	<b>92.58</b>	92.54	-0.04	<b>68.43</b>	<b>68.39</b>	-0.04	<b>67.31</b>	<b>67.30</b>	-0.01

# Yes! Better Diffusion Models are Indeed Better

- Alleviate overfitting in adversarial training

Generated	Best epoch	Clean			PGD-40			AA		
		Best	Last	Diff	Best	Last	Diff	Best	Last	Diff
x	91	84.55	82.59	-1.96	55.66	46.47	-9.19	54.37	45.29	-9.08
50K	171	86.15	85.47	-0.68	56.96	50.02	-6.94	55.71	48.85	-6.86
100K	274	88.20	87.47	-0.73	59.85	54.95	-4.90	58.85	53.42	-5.43
200K	365	89.71	89.48	-0.23	61.69	60.32	-1.37	59.91	59.11	-0.80
500K	395	90.76	90.58	-0.18	63.85	63.69	-0.16	62.76	62.77	+0.01
1M	394	91.13	90.89	-0.24	64.67	64.50	-0.17	63.35	63.50	+0.15
5M	395	91.15	90.93	-0.22	64.88	64.88	0	64.05	64.05	0
10M	396	<b>91.25</b>	<b>91.18</b>	-0.07	65.03	64.96	-0.07	64.19	64.28	+0.09
20M	399	91.17	91.07	-0.10	65.21	65.13	-0.08	64.27	64.16	-0.11
50M	395	91.24	91.15	-0.09	<b>65.35</b>	<b>65.23</b>	-0.12	<b>64.53</b>	<b>64.51</b>	-0.02

# Yes! Better Diffusion Models are Indeed Better

	Step	FID ↓	Clean	PGD-40	AA
Class-cond.	5	35.54	88.92	57.33	57.78
	10	2.477	90.96	66.21	62.81
	15	1.848	91.05	64.56	63.24
	20	<b>1.824</b>	<b>91.12</b>	<b>64.61</b>	<b>63.35</b>
	25	1.843	91.07	64.59	63.31
	30	1.861	91.10	64.51	63.25
	35	1.874	91.01	64.55	63.13
	40	1.883	91.03	64.44	63.03
	5	37.78	88.00	56.92	57.19
Uncond.	10	2.637	89.40	62.88	61.92
	15	1.998	89.36	63.47	62.31
	20	<b>1.963</b>	<b>89.76</b>	<b>63.66</b>	<b>62.45</b>
	25	1.977	89.61	63.63	62.40
	30	1.992	89.52	63.51	62.33
	35	2.003	89.39	63.56	62.37
	40	2.011	89.44	63.30	62.24

- Conditional > Unconditional
- Lower FID is better

# Yes! Better Diffusion Models are Indeed Better

**Table 6.** Test accuracy (%) with different **augmentation methods** under the  $(\ell_\infty, \epsilon = 8/255)$  threat model on CIFAR-10, using WRN-28-10 and 1M EDM generated data.

Methed	Clean	PGD-40	AA
Common	91.12	<b>64.61</b>	<b>63.35</b>
Cutout	<b>91.25</b>	64.54	63.30
CutMix	91.08	64.34	62.81
AutoAugment	91.23	64.07	62.86
RandAugment	91.14	64.39	63.12
IDBH	91.08	64.41	63.24

- Data augmentation seems ineffective

# Future Research

## Trustworthy Diffusion Models (or LLMs):

- Practical and intuitive definitions on untrustworthiness  
(trustworthiness stems from social need, do not stick to elegant math)
- Scalable tools for evaluating untrustworthiness  
(e.g., *training data extraction as a scalable way to measure privacy*)
- Finding ways to alleviate untrustworthiness  
(*find bugs and fix bugs*)

## Diffusion Models for Trustworthy ML:

- Adversarial training (high training cost)
- Adversarial purification (high inference cost)
- How to more efficiently exploit diffusion models?

# Acknowledgement



**Yunqing Zhao**



**Zekai Wang**



**Tianyu Pang**



**Chao Du**



**Min Lin**



**Shuicheng Yan**

# References



[1] A Recipe for Watermarking Diffusion Models. arXiv:2303.10137

<https://github.com/yunqing-me/WatermarkDM>

[2] Better Diffusion Models Further Improve Adversarial Training. arXiv:2302.04638

<https://github.com/wzekai99/dm-improves-at>

[3] Bag of Tricks for Training Data Extraction from Language Models. arXiv:2302.04460

[Code to be released](#)